

# Logistic Regression

## Import necessary packages

```
import numpy as np
import matplotlib.pyplot as plt
import h5py # common package to interact with a dataset that is stored
on an H5 file.
import scipy
# from PIL import Image
from scipy import ndimage
```

## Problem statement

Dataset ("data.h5") contains:

- a training set of `m_train` images labeled as cat (`y=1`) or non-cat (`y=0`)
- a test set of `m_test` images labeled as cat or non-cat
- each image is of shape `(num_px, num_px, 3)` where 3 is for the 3 channels (RGB) (each image is square: height = `num_px` and width = `num_px`).

Task is to build a simple image-classification algorithm that can correctly classify pictures as cat or non-cat.

## Overview dataset

### Create shortcut for path

```
import os
cwd= os.getcwd() # current working directory
path = os.path.join(cwd, 'data')
# print (path)
```

## Load dataset

```
def load_dataset():
    file_name= os.path.join(path , 'train_catvnoncat.h5')
    train_dataset = h5py.File(file_name, "r")
    X_train = np.array(train_dataset["train_set_x"][:]) # your train
set features
    Y_train = np.array(train_dataset["train_set_y"][:]) # your train
set labels

    file_name= os.path.join(path , 'test_catvnoncat.h5')
    test_dataset = h5py.File(file_name, "r")
    X_test = np.array(test_dataset["test_set_x"][:]) # your test set
features
    Y_test = np.array(test_dataset["test_set_y"][:]) # your test set
labels

    classes = ['non-cat', 'cat']

    Y_train = Y_train.reshape(-1,1)
    Y_test = Y_test.reshape(-1,1)

    return X_train, Y_train, X_test, Y_test, classes

X_train,Y_train, X_test, Y_test, classes = load_dataset()
# Note : in case file not found, uncomment to print path in previous
step and correct to necessary extension

print ('X_train.shape= ',X_train.shape)
print ('X_test.shape= ',X_test.shape)
print ('Y_train.shape= ',Y_train.shape)
print ('Y_test.shape= ',Y_test.shape)

X_train.shape= (209, 64, 64, 3)
X_test.shape= (50, 64, 64, 3)
Y_train.shape= (209, 1)
Y_test.shape= (50, 1)

# YOUR_CODE. get m_train, num_px and m_test
# START_CODE
m_train = X_train.shape[0]
```

```
num_px = X_train.shape[1]
m_test = X_test.shape[0]
# END_CODE
```

## Check result

```
print ("Number of training examples: m_train = " + str(m_train))
print ("Number of testing examples: m_test = " + str(m_test))
print ("Height/Width of each image: num_px = " + str(num_px))
print ("Each image is of size: (" + str(num_px) + ", " + str(num_px) +
", 3)")
```

```
Number of training examples: m_train = 209
Number of testing examples: m_test = 50
Height/Width of each image: num_px = 64
Each image is of size: (64, 64, 3)
```


## Expected output

```
Number of training examples: m_train = 209 Number of testing examples:
m_test = 50 Height/Width of each image: num_px = 64 Each image is of
size: (64, 64, 3)
```

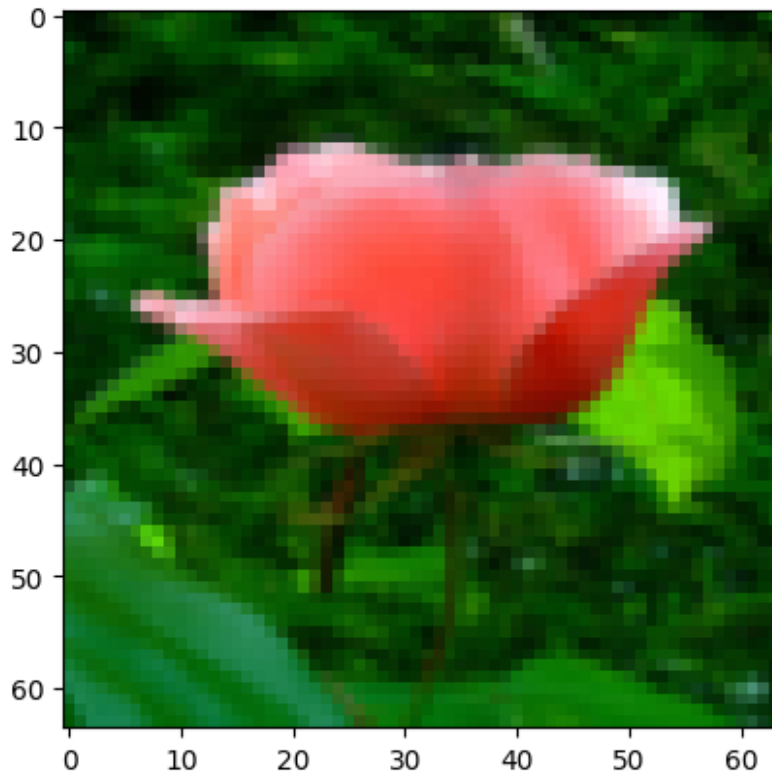
## Review some examples

```
# Example of a picture
index = 11
plt.imshow(X_train[index])
print ("y = " + str(Y_train[index,:]) + ", it's a '" +
classes[np.squeeze(Y_train[index,:])] + "' picture.")

y = [1], it's a 'cat' picture.
```



```
index = 3
plt.imshow(X_train[index])
print ("y = " + str(Y_train[index,:]) + ", it's a '" +
classes[np.squeeze(Y_train[index,:])] + "' picture.")
y = [0], it's a 'non-cat' picture.
```



```
print ('Label 1 count:', np.sum(Y_train!=0))  
print ('Label 0 count:', np.sum(Y_train==0))
```

```
Label 1 count: 72  
Label 0 count: 137
```

## Flatten features

```
# YOUR_CODE. Reshape the training and test set to shape  
(number_of_samples, num_px*num_px*3)  
# START_CODE  
X_train_flatten = X_train.reshape(m_train, -1)  
X_test_flatten = X_test.reshape(m_test, -1)  
# END_CODE
```

## Check result

```

print ("train_set_x_flatten shape: {}".format(X_train_flatten.shape))
print ("test_set_x_flatten shape: {}".format(X_test_flatten.shape))
print ("sanity check after reshaping:
{}".format(X_train_flatten[0, :5]))

train_set_x_flatten shape: (209, 12288)
test_set_x_flatten shape: (50, 12288)
sanity check after reshaping: [17 31 56 22 33]

```

## Expected Output

```

X_train_flattern shape: (209,12288) X_test_flattern shape: (50,12288)
sanity check after reshaping: [17 31 56 22 33]

```

## Normalize features

To represent color images, the red, green and blue channels (RGB) must be specified for each pixel, and so the pixel value is actually a vector of three numbers ranging from 0 to 255.

For picture datasets, it is almost the same as MinMaxScaler to just divide every row of the dataset by 255 (the maximum value of a pixel channel).

```

X_train_scaled = X_train_flatten/255.
X_test_scaled = X_test_flatten/255.

```

## sigmoid function

$$g(z) = \frac{1}{1 + e^{-z}}$$

```

def sigmoid(z):
    """
    Compute the sigmoid of z

    Arguments:
    z -- A scalar or numpy array of any size.

    Return:
    g -- sigmoid(z)
    """

```

```

# YOUR_CODE. Implement sigmoid function
# START_CODE
g = 1/(1+np.exp(-z))
# END_CODE

return g

```

Check result

```

print ("sigmoid([0, 2]) = " + str(sigmoid(np.array([0,2]))))
sigmoid([0, 2]) = [0.5          0.88079708]

```

Expected Output

```
sigmoid([0, 2]) [0.5          0.88079708]
```

## Initialize parameters

```

def initialize_with_zeros(dim):
    """
    This function creates a vector of zeros of shape (1,dim) for w and
    initializes b to 0.

    Argument:
    dim -- size of the w vector we want (or number of parameters in
    this case)

    Returns:
    w -- initialized vector of shape (1,dim)
    b -- initialized scalar (corresponds to the bias)
    """

    # YOUR_CODE. Initialize b to zero and w as a vector of zeros.
    # START_CODE
    w = np.zeros((1, dim))

```

```

b = 0
# END_CODE

assert(w.shape == (1,dim))
assert(isinstance(b, float) or isinstance(b, int))

return w, b

```

## Check result

```

dim = 2
w, b = initialize_with_zeros(dim)
print ("w = " + str(w))
print ("b = " + str(b))

w = [[0. 0.]]
b = 0

```

## Expected Output

```

w = [[ 0.  0.]]
b = 0

```

**Note:** For image inputs, w will be of shape (1, num\_px × num\_px × 3).

## Forward and Backward propagation

Computing cost function is called "forward" and computing derivatives (gradient) is called "backward" propagation

### Forward Propagation:

- compute  $Z = b + X @ w . T \left( \left[ \left[ z^{(0)} \right], \left[ z^{(1)} \right], \dots, \left[ z^{(m-1)} \right] \right] \right)$
- compute  $A = g(Z) \left( \left[ \left[ a^{(0)} \right], \left[ a^{(1)} \right], \dots, \left[ a^{(m-1)} \right] \right] \right)$
- calculate the cost function:  $J = -\frac{1}{m} \sum_{i=0}^{m-1} \left( y^{(i)} \log(a^{(i)}) + (1 - Y^{(i)}) \log(1 - a^{(i)}) \right)$



## Backward Propagation:

$$\frac{\partial J}{\partial w} = \frac{1}{m} (A - Y)^T @ X \quad \frac{\partial J}{\partial b} = \frac{1}{m} \sum_{i=1}^m (a^{(i)} - y^{(i)})$$

```
# GRADED FUNCTION: propagate
```

```
def propagate(w, b, X, Y, C=1):  
    """
```

```
    Implement the cost function and its gradient for the propagation  
    explained above
```

```
    Arguments:
```

```
    w -- weights, a numpy array of size (1,num_px * num_px * 3)
```

```
    b -- bias, a scalar
```

```
    X -- data of size (number of examples, num_px * num_px * 3)
```

```
    Y -- true "label" vector (containing 0 if non-cat, 1 if cat) of  
    size (number of examples,1)
```

```
    Return:
```

```
    cost -- negative log-likelihood cost for logistic regression
```

```
    dw -- gradient of the loss with respect to w, thus same shape as w
```

```
    db -- gradient of the loss with respect to b, thus same shape as b
```

```
    """
```

```
    m = X.shape[0]
```

```
    # YOUR_CODE. implement forward propagation
```

```
    # START_CODE
```

```
    Z = b + X @ w.T
```

```
    A = sigmoid(Z)
```

```
    cost = (-1 / m) * np.sum(Y * np.log(A) + (1 - Y) * np.log(1 - A))  
+ 1 / (2 * m) * np.sum(w ** 2)
```

```
    # END_CODE
```

```
    # YOUR_CODE. Implement Backward propagation
```

```
    # START_CODE
```

```
    dJ_dw = (1 / m) * (A - Y).T @ X + 1 / m * w
```

```
    dJ_db = (1 / m) * np.sum(A - Y)
```

```
    # END_CODE
```

```
    assert (dJ_dw.shape == w.shape)
```

```
    assert (dJ_db.dtype == float)
```

```
    assert (cost.dtype == float)
```

```
    grads = {"dJ_dw": dJ_dw,  
             "dJ_db": dJ_db}
```

```
return grads, cost
```

## Check result

```
w, b, X, Y = np.array([[1., 2.]]), 2., np.array([[1.,2.,-1.],[3.,4.,-3.2]]).T, np.array([[1,0,1]]).T
grads, cost = propagate(w, b, X, Y)
print ("dJ_dw = " + str(grads["dJ_dw"]))
print ("dJ_db = " + str(grads["dJ_db"]))
print ("cost = " + str(cost))

dJ_dw = [[1.33178935 3.06173906]]
dJ_db = 0.001455578136784208
cost = 6.6348786527278865
```

## Expected Output

```
dJ_dw = [[1.33178935 3.06173906]]
dJ_db = 0.001455578136784208
cost = 6.6348786527278865
```

## Optimization

```
def optimize(w, b, X, Y, num_iterations, learning_rate, C= 1, verbose
= False):
    """
    This function optimizes w and b by running a gradient descent
    algorithm

    Arguments:
    w -- weights, a numpy array of size (1,num_px * num_px * 3)
    b -- bias, a scalar
    X -- data of size (number of examples, num_px * num_px * 3)
    Y -- true "label" vector (containing 0 if non-cat, 1 if cat) of
    size (number of examples,1)
    num_iterations -- number of iterations of the optimization loop
```

```
learning_rate -- learning rate of the gradient descent update rule
print_cost -- True to print the loss every 100 steps
```

Returns:

```
params -- dictionary containing the weights w and bias b
grads -- dictionary containing the gradients of the weights and
bias with respect to the cost function
costs -- list of all the costs computed during the optimization,
this will be used to plot the learning curve.
"""
```

```
costs = [] # keep history for plotting if necessary
```

```
for i in range(num_iterations):
```

```
    # YOUR_CODE. Call to compute cost and gradient
```

```
    # START_CODE
```

```
    grads, cost = propagate(w, b, X, Y, C=1)
```

```
    # END_CODE
```

```
    # Retrieve derivatives from grads
```

```
    dJ_dw = grads["dJ_dw"]
```

```
    dJ_db = grads["dJ_db"]
```

```
    # YOUR_CODE. Update parameters
```

```
    # START_CODE
```

```
    w = w - learning_rate * dJ_dw
```

```
    b = b - learning_rate * dJ_db
```

```
    # END_CODE
```

```
    # Record the costs
```

```
    if i % 100 == 0:
```

```
        costs.append(cost)
```

```
    # Print the cost every 100 training iterations
```

```
    if verbose and i % 100 == 0:
```

```
        print ("Cost after iteration %i: %f" % (i, cost))
```

```
params = {"w": w,
          "b": b}
```

```
grads = {"dJ_dw": dJ_dw,
         "dJ_db": dJ_db}
```

```
return params, grads, costs
```

## Check result

```
params, grads, costs = optimize(w, b, X, Y, num_iterations= 100,
learning_rate = 0.009, verbose = False)

print ("w = " + str(params["w"]))
print ("b = " + str(params["b"]))
print ("dw = " + str(grads["dJ_dw"]))
print ("db = " + str(grads["dJ_db"]))

w = [[ 0.08006006 -0.02399336]]
b = 1.9060971483059892
dw = [[0.62090316 1.19256883]]
db = 0.2084129285706479
```

## Expected Output

```
w = [[ 0.08006006 -0.02399336]]
b = 1.9060971483059892
dw = [[0.62090316 1.19256883]]
db = 0.2084129285706479
```

## Predict

1. Calculate  $\hat{Y} = A = g(b + X @ w.T)$
2. Convert the entries of a into 0 (if activation  $\leq 0.5$ ) or 1 (if activation  $> 0.5$ ), store the predictions in a vector `Y_prediction`. Try to avoid `for` loop but use vectorized way if possible.

```
def predict(w, b, X):
    """
    Predict whether the label is 0 or 1 using learned logistic
    regression parameters (w, b)

    Arguments:
    w - weights, a numpy array of size (1,num_px * num_px * 3)
    b - bias, a scalar
    X - data of size (number of examples, num_px * num_px * 3)
```

```

Returns:
Y_prediction - a numpy array of shape (number of examples, 1)
containing all predictions (0/1) for the examples in X
'''
m,n = X.shape
assert (w.shape==(1,n))

# YOUR_CODE. Compute "A" predicting the probabilities of a cat
being present in the picture
# START_CODE
A= sigmoid(b + X @ w.T)
# END_CODE

# YOUR_CODE. Convert probabilities to actual predictions 0 or 1
# START_CODE
Y_prediction = (A >= 0.5).astype(int)
# END_CODE

assert(Y_prediction.shape == (m, 1))

return Y_prediction

```

Check result

```

w = np.array([[0.1124579],[0.23106775]]).T
b = -0.3
X = np.array([[1.,-1.1,-3.2],[1.2,2.,0.1]]).T
print ("predictions = \n{}".format (predict(w, b, X)))

predictions =
[[1]
 [1]
 [0]]

```

Expected Output

```
predictions= [[1] [1] [0]]
```

# Model

```
def model(X_train, Y_train, X_test, Y_test, num_iterations = 2000,
learning_rate = 0.5, verbose = False, C= 1):
    """
    Builds the logistic regression model by calling the functions
    implemented previously

    Arguments:
    X_train -- training set represented by a numpy array of shape
    (number of examples, num_px * num_px * 3)
    Y_train -- training labels represented by a numpy array (vector)
    of shape (1, m_train)
    X_test -- test set represented by a numpy array of shape (num_px *
    num_px * 3, m_test)
    Y_test -- test labels represented by a numpy array (vector) of
    shape (number of examples,1)
    num_iterations -- hyperparameter representing the number of
    iterations to optimize the parameters
    learning_rate -- hyperparameter representing the learning rate
    used in the update rule of optimize()
    print_cost -- Set to true to print the cost every 100 iterations
    C- regularization parameter

    Returns:
    res -- dictionary containing information about the model.
    """

    # YOUR_CODE.
    # START_CODE

    # initialize parameters
    dim = X_train.shape[1]
    w, b = initialize_with_zeros(dim)

    # run gradient descent
    parameters, grads, costs = optimize(w, b, X_train, Y_train,
num_iterations, learning_rate, C= 1, verbose=verbose)

    # retrieve parameters w and b from dictionary "parameters"
    w = parameters.get('w')
    b = parameters.get('b')

    # predict test/train set examples
    Y_prediction_test = predict(w, b, X_test)
    Y_prediction_train = predict(w, b, X_train)
```

```

# END_CODE

# Print train/test Errors
print("train accuracy= {:.3%}".format(np.mean(Y_prediction_train == Y_train)))
print("test accuracy= {:.3%}".format(np.mean(Y_prediction_test == Y_test)))

res = {'costs': costs,
       'Y_prediction_test': Y_prediction_test,
       'Y_prediction_train': Y_prediction_train,
       'w' : w,
       'b' : b,
       'learning_rate' : learning_rate,
       'num_iterations': num_iterations,
       'C':C
      }

return res

```

## Check result

```

res = model(X_train= X_train_scaled,
            Y_train=Y_train,
            X_test=X_test_scaled,
            Y_test= Y_test,
            num_iterations = 3000,
            learning_rate = 0.005,
            verbose = True,
            C= 0.3 # 0.6 is still overfitting, 0.3 is low value to
keep the test accuracy as high as possible
)

```

```

Cost after iteration 0: 0.693147
Cost after iteration 100: 0.584911
Cost after iteration 200: 0.468157
Cost after iteration 300: 0.377985
Cost after iteration 400: 0.333494
Cost after iteration 500: 0.305956
Cost after iteration 600: 0.283224
Cost after iteration 700: 0.264051
Cost after iteration 800: 0.247613
Cost after iteration 900: 0.233335
Cost after iteration 1000: 0.220802
Cost after iteration 1100: 0.209705

```

```
Cost after iteration 1200: 0.199806
Cost after iteration 1300: 0.190921
Cost after iteration 1400: 0.182902
Cost after iteration 1500: 0.175630
Cost after iteration 1600: 0.169008
Cost after iteration 1700: 0.162955
Cost after iteration 1800: 0.157403

Cost after iteration 1900: 0.152296
Cost after iteration 2000: 0.147584
Cost after iteration 2100: 0.143226
Cost after iteration 2200: 0.139186
Cost after iteration 2300: 0.135433
Cost after iteration 2400: 0.131939
Cost after iteration 2500: 0.128679
Cost after iteration 2600: 0.125634
Cost after iteration 2700: 0.122783
Cost after iteration 2800: 0.120111
Cost after iteration 2900: 0.117603
train accuracy= 99.522%
test accuracy= 68.000%
```


## Expected Output

```
...
Cost after iteration 2600: 0.057748
Cost after iteration 2700: 0.056804
Cost after iteration 2800: 0.055918
Cost after iteration 2900: 0.055084
train accuracy= 98.565%
test accuracy= 70.000%
```

## Visualize cost function changes

```
costs = np.squeeze(res['costs'])
plt.plot(costs)
plt.ylabel('cost')
plt.xlabel('iterations (per hundreds)')
plt.title("Learning rate = " + str(res["learning_rate"]))
plt.show()
```






## Visualize prediction

```
Y_test[index,0], res["Y_prediction_test"][index,0]
(1, 1)
index = 1


plt.imshow(X_test[index,:].reshape(num_px, num_px, 3))
y_true = Y_test[index,0]
y_predicted = res["Y_prediction_test"][index,0]
print ('y_predicted = {} (true label = {}) , you predicted that it is
a {} picture.\n
        format(y_predicted,
               y_true,
               classes[y_predicted]))
# todo it looks like it does not use predicted value but test label
instead
```

`y_predicted = 1` (true label = 1) , you predicted that it is a cat picture.



```
# index = 6 is example of a picture that was wrongly classified.
index = 6
plt.imshow(X_test[index,:].reshape(num_px, num_px, 3))
y_true = Y_test[index,0]
y_predicted = res["Y_prediction_test"][index,0]
print ('y_predicted = {} (true label = {}) , you predicted that it is
a {} picture.'. \
        format(y_predicted,
                y_true,
                classes[y_predicted]))
# todo it looks like it does not use predicted value but test label
instead

y_predicted = 0 (true label = 1) , you predicted that it is a non-cat
picture.
```



```
Y_test[index,0]
```

```
1
```

## Test with your image

```
from PIL import Image
file_name= os.path.join(path , 'Oleksiy.Tsebriy.jpg')
image = Image.open(file_name).resize((num_px,num_px))
print ('image.size= ', image.size)
image
image.size= (64, 64)
```



```

my_image= np.array(image.getdata())
my_image.shape

(4096, 3)

my_image= my_image.reshape((1, num_px*num_px*3))
print ('my_image.shape=',my_image.shape)
my_image

my_image.shape= (1, 12288)

array([[227, 227, 217, ..., 43, 65, 96]])

import warnings
warnings.filterwarnings("ignore")

my_predicted_image = predict(res["w"], res["b"], my_image)
my_predicted_image
print('y = {} , your algorithm predicts a {} picture.'.

format(np.squeeze(my_predicted_image),classes[np.squeeze(my_predicted_
image)]))

warnings.filterwarnings("default")

y = 0 , your algorithm predicts a non-cat picture.

file_name= os.path.join(path , 'test_cat.jpg')
image = Image.open(file_name).resize((num_px,num_px))
print ('image.size= ', image.size)
image = image.convert('RGB')
image

image.size= (64, 64)

```



```

my_image= np.array(image.getdata())[:, :3] # by unknown reason picture
made as screenshot has 4 channels
print ('my_image.shape=', my_image.shape)
my_image= my_image.reshape((1, num_px*num_px*3))
print ('after reshape: my_image.shape=',my_image.shape)
my_predicted_image = predict(res["w"], res["b"], my_image)

print('y = {} , your algorithm predicts a {} picture.'.

```

```

format(np.squeeze(my_predicted_image), classes[np.squeeze(my_predicted_
image)]))

my_image.shape= (4096, 3)
after reshape: my_image.shape= (1, 12288)
y = 1 , your algorithm predicts a cat picture.

```

## Sklearn implementation

```

from sklearn.linear_model import LogisticRegression
y_train = np.squeeze(Y_train) # LogisticRegression requires 1d input for y
clf = LogisticRegression(C=0.01).fit(X_train_scaled, y_train)

print("train accuracy= {:.3%}".format(clf.score (X_train_scaled,
y_train)))
y_test = np.squeeze(Y_test)
print("test accuracy= {:.3%}".format(clf.score (X_test_scaled,
y_test)))

train accuracy= 90.909%
test accuracy= 66.000%

print('y = {} , sklearn algorithm predicts a {} picture.'.

format(np.squeeze(clf.predict(my_image)), classes[np.squeeze(clf.predic
t(my_image))]))

y = 1 , sklearn algorithm predicts a cat picture.

```

## Sklearn for breast cancer dataset

```

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
# import load_breast_cancer and get the X_cancer, y_cancer
from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer()
(X_cancer, y_cancer) = cancer.data, cancer.target

# split to train and test using random_state = 0
X_train, X_test, y_train, y_test = train_test_split(X_cancer,

```

```

y_cancer)
# train LogisticRegression classifier for max_iter= 10000
clf = LogisticRegression(max_iter=10000).fit(X_train, y_train)
print('\nBreast cancer dataset')
print('X_cancer.shape= {}'.format(X_cancer.shape))
print('Accuracy of Logistic regression classifier on training set:
{:.2f}'
      .format(clf.score(X_train, y_train)))
print('Accuracy of Logistic regression classifier on test set: {:.2f}'
      .format(clf.score(X_test, y_test)))

```

```

Breast cancer dataset
X_cancer.shape= (569, 30)
Accuracy of Logistic regression classifier on training set: 0.96
Accuracy of Logistic regression classifier on test set: 0.96

```

## Expected Output

```

Breast cancer dataset
X_cancer.shape= (569, 30)
Accuracy of Logistic regression classifier on training set: 0.96
Accuracy of Logistic regression classifier on test set: 0.95

```

## Sklearn for synthetic dataset

### Additional functions for visualization

```

%matplotlib notebook

def plot_decision_boundary(clf, X_train, y_train, X_test=None, y_test=
None, title=None, precision=0.01, plot_symbol_size = 50, ax= None,
is_extended=True):
    ...
    Draws the binary decision boundary for X that is nor required
    additional features and transformation (like polynomial)

```

```

...
# Create color maps - required by pcolormesh
from matplotlib.colors import ListedColormap
colors_for_points = np.array(['grey', 'orange']) # neg/pos
colors_for_areas = np.array(['grey', 'orange']) # neg/pos # alpha
is_applied_later
cmap_light = ListedColormap(colors_for_areas)

mesh_step_size = precision #.01 # step size in the mesh
if X_test is None or y_test is None:
    show_test= False
    X= X_train
else:
    show_test= True
    X= np.concatenate([X_train,X_test], axis=0)
x1_min, x1_max = X[:, 0].min() - .1, X[:, 0].max() + 0.1
x2_min, x2_max = X[:, 1].min() - .1, X[:, 1].max() + 0.1
# Create grids of pairs
xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, mesh_step_size),
                        np.arange(x2_min, x2_max, mesh_step_size))
# Flatten all samples
target_samples_grid= (np.c_[xx1.ravel(), xx2.ravel()])

print ('Call prediction for all grid values (precision of drawing
= {},\n you may configure to speed up e.g.
precision=0.05)'.format(precision))
Z = clf.predict(target_samples_grid)

# Reshape the result to original meshgrid shape
Z = Z.reshape(xx1.shape)

if ax:
    plt.sca(ax)

# Plot all meshgrid prediction
plt.pcolormesh(xx1, xx2,Z, cmap = cmap_light, alpha=0.2)

# Plot train set
plt.scatter(X_train[:, 0], X_train[:, 1], s=plot_symbol_size,
            c=colors_for_points[y_train.ravel()], edgecolor =
'black',alpha=0.6)
# Plot test set
if show_test:
    plt.scatter(X_test[:, 0], X_test[:, 1], marker='^',
s=plot_symbol_size,
            c=colors_for_points[y_test.ravel()],edgecolor =
'black',alpha=0.6)

```

```

    if is_extended:
        # Create legend
        import matplotlib.patches as mpatches # use to assign levels
        for colored points
            patch0 = mpatches.Patch(color=colors_for_points[0],
label='negative')
            patch1 = mpatches.Patch(color=colors_for_points[1],
label='positive')
            plt.legend(handles=[patch0, patch1])
        plt.title(title)
    if is_extended:
        plt.xlabel('feature 1')
        plt.ylabel('feature 2')
    else:
        plt.tick_params(
            top =False,
            bottom= False,
            left  = False,
            labelleft = False,
            labelbottom = False
        )

def plot_data_logistic_regression(X,y,legend_loc= 1, title= None):
    """
    :param X: 2 dimensional ndarray
    :param y: 1 dimensional ndarray. Use y.ravel() if necessary
    :return:
    """

    positive_indices = (y == 1)
    negative_indices = (y == 0)
    # import matplotlib as mpl
    colors_for_points = ['grey', 'orange'] # neg/pos

    plt.scatter(X[negative_indices][:,0], X[negative_indices][:,1],
s=40, c=colors_for_points [0], edgecolor = 'black', label='negative',
alpha = 0.7)
    plt.scatter(X[positive_indices][:,0], X[positive_indices][:,1],
s=40, c=colors_for_points [1], edgecolor = 'black', label='positive',
alpha = 0.7)
    plt.title(title)
    plt.legend(loc= legend_loc)

```

Make classification




```

from sklearn.datasets import make_classification
Xc_2, yc_2= make_classification(n_samples=100,
                                n_features=2,
                                n_informative=2,
                                n_redundant=0,
                                random_state=2018,
                                n_clusters_per_class=1,
                                flip_y = 0.1,
                                class_sep = 0.8)

plt.figure()
plot_data_logistic_regression(Xc_2, yc_2, title=
'make_classification')

```



```

C = 1000
X_train, X_test, y_train, y_test = train_test_split(Xc_2, yc_2,
                                                    random_state = 0)
clf = LogisticRegression(C=C).fit(X_train, y_train)
print('Make Regression')
print('Xc_2.shape= {}'.format(Xc_2.shape))
print('Accuracy of Logistic regression classifier on training set:
{:.2f}'
      .format(clf.score(X_train, y_train)))

```

```
print('Accuracy of Logistic regression classifier on test set: {:.2f}'
      .format(clf.score(X_test, y_test)))
```

Make Regression

```
Xc_2.shape= (100, 2)
```

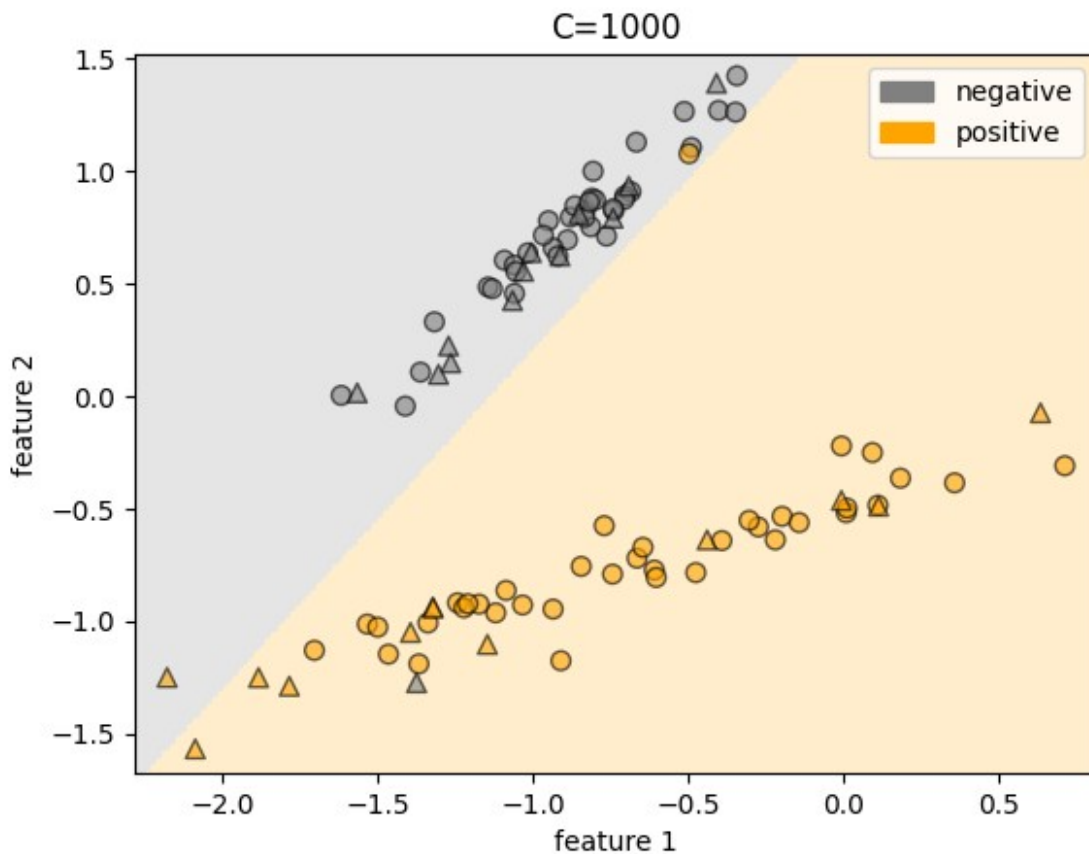
```
Accuracy of Logistic regression classifier on training set: 0.99
```

```
Accuracy of Logistic regression classifier on test set: 0.92
```

```
plt.figure()
```

```
plot_decision_boundary(clf, X_train, y_train, X_test, y_test,
title='C={}'.format(C),precision=0.01, plot_symbol_size = 50)
```

Call prediction for all grid values (precision of drawing = 0.01,  
you may configure to speed up e.g. precision=0.05)



```
_, ((ax1, ax2, ax3), (ax4, ax5, ax6)) = plt.subplots(2, 3)
axes = (ax1, ax2, ax3, ax4, ax5, ax6)
C = (0.001, 0.01, 0.1, 1, 100, 1000)
for i in range(len(C)):
    clf = LogisticRegression(C=C[i]).fit(X_train, y_train)
    print('Accuracy = {:.2f}'.format(clf.score(X_train, y_train)))

    plot_decision_boundary(clf, X_train, y_train,
```

```
title='C={}'.format(C[i]),precision=0.01, plot_symbol_size = 30, ax=
axes[i], is_extended=False)
```

Accuracy = 0.57

Call prediction for all grid values (precision of drawing = 0.01,  
you may configure to speed up e.g. precision=0.05)

Accuracy = 0.97

Call prediction for all grid values (precision of drawing = 0.01,  
you may configure to speed up e.g. precision=0.05)

Accuracy = 0.99

Call prediction for all grid values (precision of drawing = 0.01,  
you may configure to speed up e.g. precision=0.05)

Accuracy = 0.99

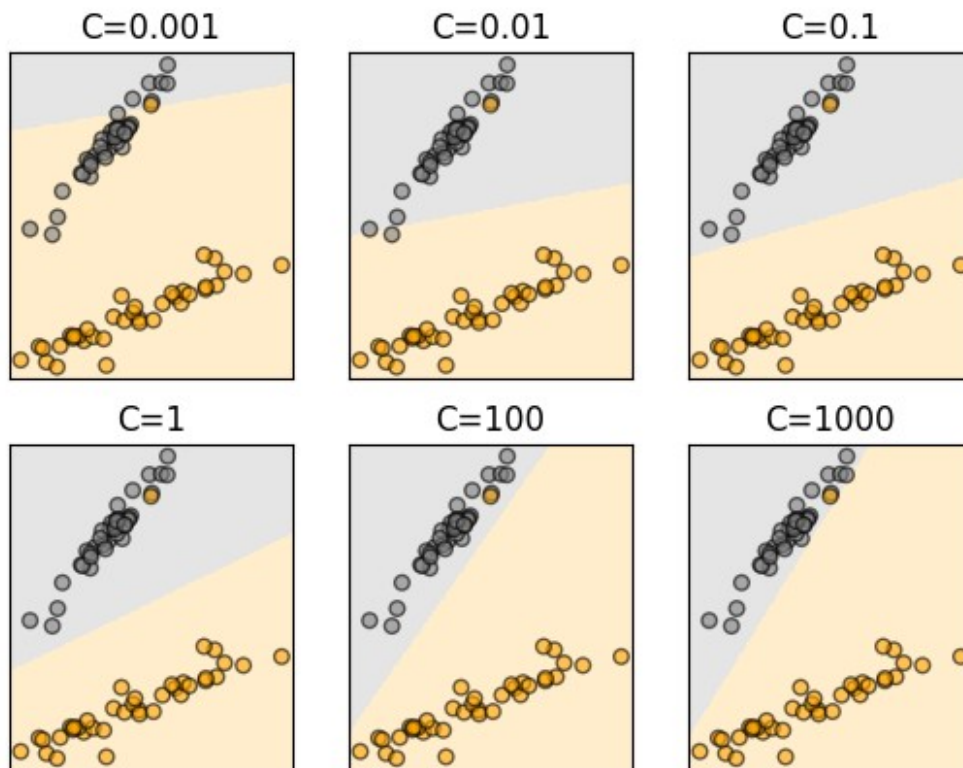
Call prediction for all grid values (precision of drawing = 0.01,  
you may configure to speed up e.g. precision=0.05)

Accuracy = 0.99

Call prediction for all grid values (precision of drawing = 0.01,  
you may configure to speed up e.g. precision=0.05)

Accuracy = 0.99

Call prediction for all grid values (precision of drawing = 0.01,  
you may configure to speed up e.g. precision=0.05)



# Polynomial In Logistic Regression

```
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
```

## Additional functions for visualization

```
%matplotlib notebook

def plot_decision_boundary(clf, X_train, y_train, X_test=None, y_test=
None, title=None, precision=0.05, plot_symbol_size = 50, ax= None,
is_extended=True):
    """
    Draws the binary decision boundary for X that is nor required
    additional features and transformation (like polynomial)
    """
    # Create color maps - required by pcolormesh
    from matplotlib.colors import ListedColormap
    colors_for_points = np.array(['grey', 'orange']) # neg/pos
    colors_for_areas = np.array(['grey', 'orange']) # neg/pos # alpha
    is applied later
    cmap_light = ListedColormap(colors_for_areas)

    mesh_step_size = precision #.01 # step size in the mesh
    if X_test is None or y_test is None:
        show_test= False
        X= X_train
    else:
        show_test= True
        X= np.concatenate([X_train,X_test], axis=0)
    x1_min, x1_max = X[:, 0].min() - .1, X[:, 0].max() + 0.1
    x2_min, x2_max = X[:, 1].min() - .1, X[:, 1].max() + 0.1
    # Create grids of pairs
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, mesh_step_size),
                           np.arange(x2_min, x2_max, mesh_step_size))
    # Flatten all samples
    target_samples_grid= (np.c_[xx1.ravel(), xx2.ravel()])
```

```

    print ('Call prediction for all grid values (precision of drawing
= {},\n you may configure to speed up e.g.
precision=0.05)'.format(precision))
    Z = clf.predict(target_samples_grid)

    # Reshape the result to original meshgrid shape
    Z = Z.reshape(xx1.shape)

    if ax:
        plt.sca(ax)

    # Plot all meshgrid prediction
    plt.pcolormesh(xx1, xx2, Z, cmap = cmap_light, alpha=0.2)

    # Plot train set
    plt.scatter(X_train[:, 0], X_train[:, 1], s=plot_symbol_size,
                c=colors_for_points[y_train.ravel()], edgecolor =
'black', alpha=0.6)
    # Plot test set
    if show_test:
        plt.scatter(X_test[:, 0], X_test[:, 1], marker='^',
s=plot_symbol_size,
                c=colors_for_points[y_test.ravel()], edgecolor =
'black', alpha=0.6)
    if is_extended:
        # Create legend
        import matplotlib.patches as mpatches # use to assign levels
        for colored points
        patch0 = mpatches.Patch(color=colors_for_points[0],
label='negative')
        patch1 = mpatches.Patch(color=colors_for_points[1],
label='positive')
        plt.legend(handles=[patch0, patch1])
    plt.title(title)
    if is_extended:
        plt.xlabel('feature 1')
        plt.ylabel('feature 2')
    else:
        plt.tick_params(
            top = False,
            bottom = False,
            left = False,
            labelleft = False,
            labelbottom = False
        )

def plot_decision_boundary_poly(clf, X_train, y_train, degree,
X_test=None, y_test= None, title=None, precision=0.05, plot_symbol_size
= 50, ax= None, is_extended=True):

```

```

'''
    Draws the binary decision boundary for X that is nor required
    additional features and transformation (like polynomial)
'''
from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree=degree,include_bias=False)

# Create color maps - required by pcolormesh
from matplotlib.colors import ListedColormap
colors_for_points = np.array(['grey', 'orange']) # neg/pos
colors_for_areas = np.array(['grey', 'orange']) # neg/pos # alpha
is applied later
cmap_light = ListedColormap(colors_for_areas)

mesh_step_size = precision #.01 # step size in the mesh
if X_test is None or y_test is None:
    show_test= False
    X= X_train
else:
    show_test= True
    X= np.concatenate([X_train,X_test], axis=0)
x1_min, x1_max = X[:, 0].min() - .1, X[:, 0].max() + 0.1
x2_min, x2_max = X[:, 1].min() - .1, X[:, 1].max() + 0.1
# Create grids of pairs
xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, mesh_step_size),
                        np.arange(x2_min, x2_max, mesh_step_size))
# Flatten all samples
target_samples_grid= (np.c_[xx1.ravel(), xx2.ravel()])
target_samples_grid_poly = poly.fit_transform(target_samples_grid)
print ('Call prediction for all grid values (precision of drawing
= {},\n you may configure to speed up e.g.
precision=0.05)'.format(precision))
Z = clf.predict(target_samples_grid_poly)
print ('Computing prediction completed.')
# Reshape the result to original meshgrid shape
Z = Z.reshape(xx1.shape)

if ax:
    plt.sca(ax)

# Plot all meshgrid prediction
plt.pcolormesh(xx1, xx2,Z, cmap = cmap_light, alpha=0.2)

# Plot train set
plt.scatter(X_train[:, 0], X_train[:, 1], s=plot_symbol_size,
            c=colors_for_points[y_train.ravel()], edgecolor =

```

```

'black',alpha=0.6)
    # Plot test set
    if show_test:
        plt.scatter(X_test[:, 0], X_test[:, 1], marker='^',
s=plot_symbol_size,
                    c=colors_for_points[y_test.ravel()],edgecolor =
'black',alpha=0.6)
    if is_extended:
        # Create legend
        import matplotlib.patches as mpatches # use to assign levels
for colored points
        patch0 = mpatches.Patch(color=colors_for_points[0],
label='negative')
        patch1 = mpatches.Patch(color=colors_for_points[1],
label='positive')
        plt.legend(handles=[patch0, patch1])
    plt.title(title)
    if is_extended:
        plt.xlabel('feature 1')
        plt.ylabel('feature 2')
    else:
        plt.tick_params(
            top=False,
            bottom=False,
            left=False,
            labelleft=False,
            labelbottom=False
        )

def plot_data_logistic_regression(X,y,legend_loc= None, title= None):
    """
    :param X: 2 dimensional ndarray
    :param y: 1 dimensional ndarray. Use y.ravel() if necessary
    :return:
    """

    positive_indices = (y == 1)
    negative_indices = (y == 0)
    # import matplotlib as mpl
    colors_for_points = ['grey', 'orange'] # neg/pos

    plt.scatter(X[negative_indices][:,0], X[negative_indices][:,1],
s=40, c=colors_for_points [0], edgecolor = 'black', label='negative',
alpha = 0.7)
    plt.scatter(X[positive_indices][:,0], X[positive_indices][:,1],
s=40, c=colors_for_points [1], edgecolor = 'black',label='positive',
alpha = 0.7)

```

```

plt.title(title)
plt.legend(loc= legend_loc)

def plot_multi_class_logistic_regression(X,y,dict_names=None, colors=
None, title =None):
    '''
    Draw the multi class samples of 2 features
    :param X: X 2 ndarray (m,2),
    :param y: vector (m,)
    :param dict_names: dict of values of y and names
    :return: None
    '''
    if not colors:
        colors_for_points = ['green','grey', 'orange', 'brown']
    else:
        colors_for_points = colors

    y_unique = list(set(y))

    for i in range (len(y_unique)):
        ind = y == y_unique[i] # vector

        if dict_names:
            plt.scatter(X[ind,0], X[ind,1], c=colors_for_points[i],
s=40, label=dict_names[y_unique[i]],edgecolor='black', alpha=.7)
        else:
            plt.scatter(X[ind, 0], X[ind, 1], s=40,
c=colors_for_points [i], edgecolor = 'black', alpha = 0.7)
        if title:
            plt.title(title)

    if dict_names:
        plt.legend(frameon=True)

def draw_linear_decision_boundaries_multiclass(clf,X,):
    colors= ['green','grey', 'orange', 'brown']
    x_line = np.linspace(X[:,0].min(),X[:,0].max(), 100)
    for w, b, color in zip(clf.coef_, clf.intercept_, colors):
        # Since class prediction with a linear model uses the formula
 $y = w_0 x_0 + w_1 x_1 + b$ ,
        # and the decision boundary is defined as being all points
        with  $y = 0$ , to plot  $x_1$  as a
        # function of  $x_0$  we just solve  $w_0 x_0 + w_1 x_1 + b = 0$  for
 $x_1$ :
        y_line = -(x_line * w[0] + b) / w[1]
        ind = (X[:,0].min()< x_line) & (x_line <X[:,0].max()) &
(X[:,1].min()< y_line) & (y_line <X[:,1].max() )
        plt.plot(x_line[ind] , y_line[ind], '-', c=color, alpha=.8)

from sklearn.datasets import make_blobs

```




```

X_mk8, y_mk8 = make_blobs(n_samples = 200, n_features = 2, centers =
8, # centers impacts for y
                           cluster_std = 1.3, random_state = 4)
y_train = y_mk8 % 2 # make it binary since make_blobs with centers =
8 creates y in [0..7]

plt.figure()
plot_data_logistic_regression(X_mk8,y_train)

```



## Polynomial Features

```

from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LogisticRegression

degree = 10
poly= PolynomialFeatures(degree=degree,include_bias=False) # default
is True means to return the first feature of all 1 as for degree 0
X_train_poly= poly.fit_transform(X_mk8)

C = 0.01
clf = LogisticRegression(C=C).fit(X_train_poly, y_train)

```

```
accuracy = clf.score (X_train_poly, y_train)
# print("train accuracy= {:.3%}".format(accuracy))
```

c:\Users\Smeeek\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

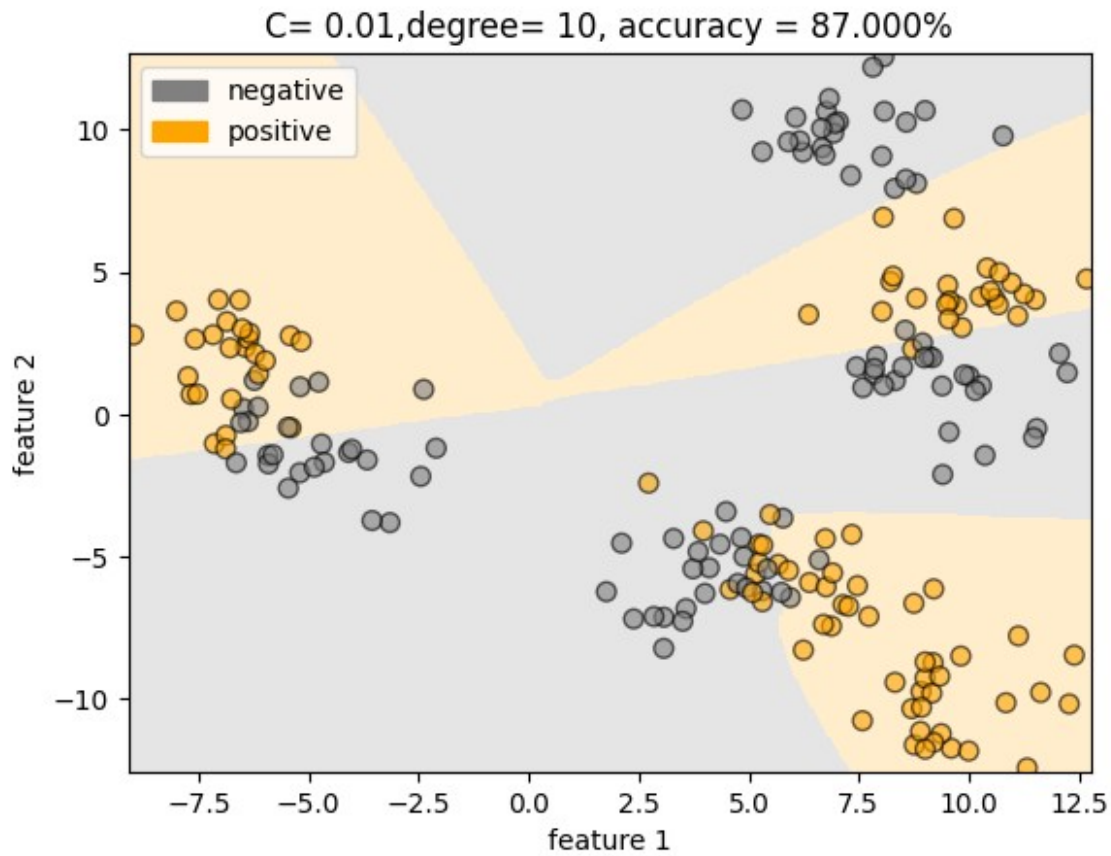
```
n_iter_i = _check_optimize_result(
```

```
plt.figure()
```

```
ax = plt.gca()
```

```
plot_decision_boundary_poly(clf, X_train_poly, y_train, degree=
degree, ax = ax, precision= 0.05, title = 'C= {},degree= {}, accuracy
= {:.3%}'.format(C, degree, accuracy))
```

Call prediction for all grid values (precision of drawing = 0.05,  
you may configure to speed up e.g. precision=0.05)  
Computing prediction completed.

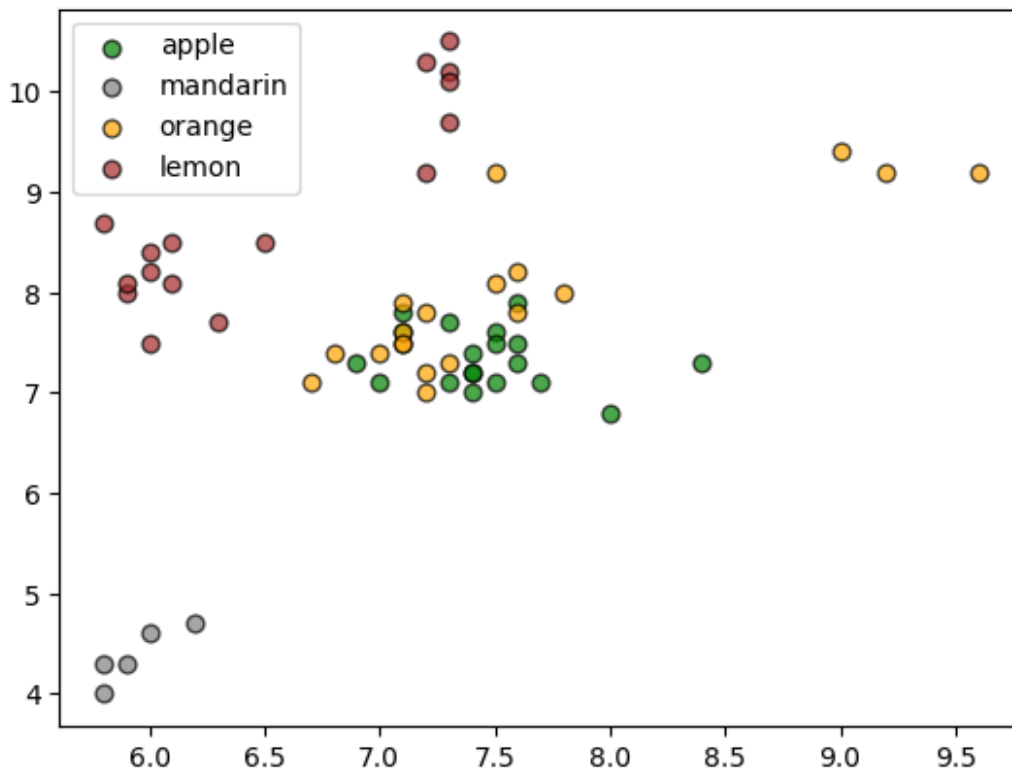


## Muticlass classification

Load fruits data set

```
import os
cwd= os.getcwd() # current working directory
path = os.path.join(cwd, 'data')
fn= os.path.join(path , 'fruit_data_with_colors.txt')
df_fruits = pd.read_table(fn)
X = df_fruits[['width', 'height']].values
y = df_fruits['fruit_label'].values
df_fruits.head(20)
fruits_dict = dict(zip(df_fruits['fruit_label'].unique(),
df_fruits['fruit_name'].unique()))
# fruits_dict
```

```
plt.figure()
plot_multi_class_logistic_regression (X,y,dict_names=fruits_dict)
```



```
help(LogisticRegression)
```

Help on class LogisticRegression in module  
sklearn.linear\_model.\_logistic:

```
class
LogisticRegression(sklearn.linear_model._base.LinearClassifierMixin,
sklearn.linear_model._base.SparseCoefMixin,
sklearn.base.BaseEstimator)
| LogisticRegression(penalty='l2', *, dual=False, tol=0.0001, C=1.0,
fit_intercept=True, intercept_scaling=1, class_weight=None,
random_state=None, solver='lbfgs', max_iter=100, multi_class='auto',
verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
|
| Logistic Regression (aka logit, MaxEnt) classifier.
|
| In the multiclass case, the training algorithm uses the one-vs-
rest (OvR)
| scheme if the 'multi_class' option is set to 'ovr', and uses the
| cross-entropy loss if the 'multi_class' option is set to
'multinomial'.
| (Currently the 'multinomial' option is supported only by the
```

```

'lbfgs',
|   'sag', 'saga' and 'newton-cg' solvers.)
|
|   This class implements regularized logistic regression using the
|   'liblinear' library, 'newton-cg', 'sag', 'saga' and 'lbfgs'
solvers. **Note
|   that regularization is applied by default**. It can handle both
dense
|   and sparse input. Use C-ordered arrays or CSR matrices containing
64-bit
|   floats for optimal performance; any other input format will be
converted
|   (and copied).
|
|   The 'newton-cg', 'sag', and 'lbfgs' solvers support only L2
regularization
|   with primal formulation, or no regularization. The 'liblinear'
solver
|   supports both L1 and L2 regularization, with a dual formulation
only for
|   the L2 penalty. The Elastic-Net regularization is only supported
by the
|   'saga' solver.
|
|   Read more in the :ref:`User Guide <logistic_regression>`.
|
|   Parameters
|   -----
|   penalty : {'l1', 'l2', 'elasticnet', None}, default='l2'
|       Specify the norm of the penalty:
|
|       - `None`: no penalty is added;
|       - `l2`: add a L2 penalty term and it is the default choice;
|       - `l1`: add a L1 penalty term;
|       - `elasticnet`: both L1 and L2 penalty terms are added.
|
|   .. warning::
|       Some penalties may not work with some solvers. See the
parameter
|       `solver` below, to know the compatibility between the
penalty and
|       solver.
|
|   .. versionadded:: 0.19
|       l1 penalty with SAGA solver (allowing 'multinomial' + L1)
|
|   .. deprecated:: 1.2
|       The 'none' option was deprecated in version 1.2, and will
be removed

```

```

    in 1.4. Use `None` instead.

    dual : bool, default=False
        Dual (constrained) or primal (regularized, see also
        :ref:`this equation <regularized-logistic-loss>`) formulation.
Dual formulation
    is only implemented for l2 penalty with liblinear solver.
Prefer dual=False when
    n_samples > n_features.

    tol : float, default=1e-4
        Tolerance for stopping criteria.

    C : float, default=1.0
        Inverse of regularization strength; must be a positive float.
        Like in support vector machines, smaller values specify
stronger
        regularization.

    fit_intercept : bool, default=True
        Specifies if a constant (a.k.a. bias or intercept) should be
        added to the decision function.

    intercept_scaling : float, default=1
        Useful only when the solver 'liblinear' is used
        and self.fit_intercept is set to True. In this case, x becomes
        [x, self.intercept_scaling],
        i.e. a "synthetic" feature with constant value equal to
        intercept_scaling is appended to the instance vector.
        The intercept becomes ``intercept_scaling *
synthetic_feature_weight``.

        Note! the synthetic feature weight is subject to l1/l2
regularization
        as all other features.
        To lessen the effect of regularization on synthetic feature
weight
        (and therefore on the intercept) intercept_scaling has to be
increased.

    class_weight : dict or 'balanced', default=None
        Weights associated with classes in the form ``{class_label:
weight}``.
        If not given, all classes are supposed to have weight one.

        The "balanced" mode uses the values of y to automatically
adjust
        weights inversely proportional to class frequencies in the
input data

```

```

    as ``n_samples / (n_classes * np.bincount(y))``.

    Note that these weights will be multiplied with sample_weight
    (passed through the fit method) if sample_weight is specified.

    .. versionadded:: 0.17
       *class_weight='balanced'*

    random_state : int, RandomState instance, default=None
        Used when ``solver`` == 'sag', 'saga' or 'liblinear' to
        shuffle the data. See :term:`Glossary <random_state>` for details.

    solver : {'lbfgs', 'liblinear', 'newton-cg', 'newton-cholesky',
              'sag', 'saga'}, default='lbfgs'
        Algorithm to use in the optimization problem. Default is
        'lbfgs'.
        To choose a solver, you might want to consider the following
        aspects:

        - For small datasets, 'liblinear' is a good choice,
        whereas 'sag'
        and 'saga' are faster for large ones;
        - For multiclass problems, only 'newton-cg', 'sag', 'saga'
        and
        'lbfgs' handle multinomial loss;
        - 'liblinear' is limited to one-versus-rest schemes.
        - 'newton-cholesky' is a good choice for `n_samples` >>
        `n_features`,
        especially with one-hot encoded categorical features
        with rare
        categories. Note that it is limited to binary
        classification and the
        one-versus-rest reduction for multiclass classification.
        Be aware that
        the memory usage of this solver has a quadratic
        dependency on
        `n_features` because it explicitly computes the Hessian
        matrix.

    .. warning::
        The choice of the algorithm depends on the penalty chosen.
        Supported penalties by solver:

        - 'lbfgs'          - ['l2', None]
        - 'liblinear'      - ['l1', 'l2']
        - 'newton-cg'      - ['l2', None]

```

```

- 'newton-cholesky' - ['l2', None]
- 'sag' - ['l2', None]
- 'saga' - ['elasticnet', 'l1', 'l2', None]

.. note::
    'sag' and 'saga' fast convergence is only guaranteed on
features
    with approximately the same scale. You can preprocess the
data with
    a scaler from :mod:`sklearn.preprocessing`.

.. seealso::
    Refer to the User Guide for more information regarding
    :class:`LogisticRegression` and more specifically the
    :ref:`Table <Logistic_regression>`
    summarizing solver/penalty supports.

.. versionadded:: 0.17
    Stochastic Average Gradient descent solver.
.. versionadded:: 0.19
    SAGA solver.
.. versionchanged:: 0.22
    The default solver changed from 'liblinear' to 'lbfgs' in
0.22.

.. versionadded:: 1.2
    newton-cholesky solver.

max_iter : int, default=100
    Maximum number of iterations taken for the solvers to
converge.

multi_class : {'auto', 'ovr', 'multinomial'}, default='auto'
    If the option chosen is 'ovr', then a binary problem is fit
for each
    label. For 'multinomial' the loss minimised is the multinomial
loss fit
    across the entire probability distribution, *even when the
data is
    binary*. 'multinomial' is unavailable when solver='liblinear'.
    'auto' selects 'ovr' if the data is binary, or if
solver='liblinear',
    and otherwise selects 'multinomial'.

.. versionadded:: 0.18
    Stochastic Average Gradient descent solver for
'multinomial' case.
.. versionchanged:: 0.22
    Default changed from 'ovr' to 'auto' in 0.22.

```



```

| verbose : int, default=0
|     For the liblinear and lbfgs solvers set verbose to any
positive
|     number for verbosity.
|
| warm_start : bool, default=False
|     When set to True, reuse the solution of the previous call to
fit as
|     initialization, otherwise, just erase the previous solution.
|     Useless for liblinear solver. See :term:`the Glossary
<warm_start>`.
|
| .. versionadded:: 0.17
|     *warm_start* to support *lbfgs*, *newton-cg*, *sag*, *saga*
solvers.
|
| n_jobs : int, default=None
|     Number of CPU cores used when parallelizing over classes if
|     multi_class='ovr'. This parameter is ignored when the
``solver`` is
|     set to 'liblinear' regardless of whether 'multi_class' is
specified or
|     not. ``None`` means 1 unless in
a :obj:`joblib.parallel_backend`
|     context. ``-1`` means using all processors.
|     See :term:`Glossary <n_jobs>` for more details.
|
| l1_ratio : float, default=None
|     The Elastic-Net mixing parameter, with ``0 <= l1_ratio <= 1``.
Only
|     used if ``penalty='elasticnet'``. Setting ``l1_ratio=0`` is
equivalent
|     to using ``penalty='l2'``, while setting ``l1_ratio=1`` is
equivalent
|     to using ``penalty='l1'``. For ``0 < l1_ratio < 1``, the
penalty is a
|     combination of L1 and L2.
|
| Attributes
| -----
|
| classes_ : ndarray of shape (n_classes, )
|     A list of class labels known to the classifier.
|
| coef_ : ndarray of shape (1, n_features) or (n_classes,
n_features)
|     Coefficient of the features in the decision function.
|
|     ``coef_`` is of shape (1, n_features) when the given problem is

```

```

binary.
|     In particular, when `multi_class='multinomial'`, `coef_`
corresponds
|     to outcome 1 (True) and `-coef_` corresponds to outcome 0
(False).
|
|     intercept_ : ndarray of shape (1,) or (n_classes,)
|     Intercept (a.k.a. bias) added to the decision function.
|
|     If `fit_intercept` is set to False, the intercept is set to
zero.
|     `intercept_` is of shape (1,) when the given problem is
binary.
|     In particular, when `multi_class='multinomial'`, `intercept_`
corresponds to outcome 1 (True) and `-intercept_` corresponds
to
|     outcome 0 (False).
|
|     n_features_in_ : int
|     Number of features seen during :term:`fit`.
|
|     .. versionadded:: 0.24
|
|     feature_names_in_ : ndarray of shape (`n_features_in_`,)
|     Names of features seen during :term:`fit`. Defined only when
`X`
|     has feature names that are all strings.
|
|     .. versionadded:: 1.0
|
|     n_iter_ : ndarray of shape (n_classes,) or (1, )
|     Actual number of iterations for all classes. If binary or
multinomial,
|     it returns only 1 element. For liblinear solver, only the
maximum
|     number of iteration across all classes is given.
|
|     .. versionchanged:: 0.20
|
|     In SciPy <= 1.0.0 the number of lbfgs iterations may
exceed
|     ``max_iter``. ``n_iter_`` will now report at most
``max_iter``.
|
|     See Also
|     -----
|     SGDClassifier : Incrementally trained logistic regression (when
given
|     the parameter ``loss="log_loss"``).

```

| LogisticRegressionCV : Logistic regression with built-in cross validation.

#### | Notes

| -----

| The underlying C implementation uses a random number generator to select features when fitting the model. It is thus not uncommon, to have slightly different results for the same input data. If that happens, try with a smaller tol parameter.

| Predict output may not match that of standalone liblinear in certain

| cases. See :ref:`differences from liblinear`  
<liblinear\_differences>`  
| in the narrative documentation.

#### | References

| -----

| L-BFGS-B -- Software for Large-scale Bound-constrained  
Optimization

| Ciyou Zhu, Richard Byrd, Jorge Nocedal and Jose Luis Morales.  
| <http://users.iems.northwestern.edu/~nocedal/lbfgsb.html>

| LIBLINEAR -- A Library for Large Linear Classification  
| <https://www.csie.ntu.edu.tw/~cjlin/liblinear/>

| SAG -- Mark Schmidt, Nicolas Le Roux, and Francis Bach  
| Minimizing Finite Sums with the Stochastic Average Gradient  
| <https://hal.inria.fr/hal-00860051/document>

| SAGA -- Defazio, A., Bach F. & Lacoste-Julien S. (2014).  
| :arxiv:`"SAGA: A Fast Incremental Gradient Method With

Support

| for Non-Strongly Convex Composite Objectives" <1407.0202>`

| Hsiang-Fu Yu, Fang-Lan Huang, Chih-Jen Lin (2011). Dual coordinate  
descent

| methods for logistic regression and maximum entropy models.  
| Machine Learning 85(1-2):41-75.  
| [https://www.csie.ntu.edu.tw/~cjlin/papers/maxent\\_dual.pdf](https://www.csie.ntu.edu.tw/~cjlin/papers/maxent_dual.pdf)

#### | Examples

| -----

| >>> from sklearn.datasets import load\_iris  
| >>> from sklearn.linear\_model import LogisticRegression  
| >>> X, y = load\_iris(return\_X\_y=True)  
| >>> clf = LogisticRegression(random\_state=0).fit(X, y)  
| >>> clf.predict(X[:2, :])

```

array([0, 0])
>>> clf.predict_proba(X[:2, :])
array([[9.8...e-01, 1.8...e-02, 1.4...e-08],
       [9.7...e-01, 2.8...e-02, ...e-08]])
>>> clf.score(X, y)
0.97...

```

Method resolution order:

```

LogisticRegression
sklearn.linear_model._base.LinearClassifierMixin
sklearn.base.ClassifierMixin
sklearn.linear_model._base.SparseCoefMixin
sklearn.base.BaseEstimator
sklearn.utils._metadata_requests._MetadataRequester
builtins.object

```

Methods defined here:

```

__init__(self, penalty='l2', *, dual=False, tol=0.0001, C=1.0,
fit_intercept=True, intercept_scaling=1, class_weight=None,
random_state=None, solver='lbfgs', max_iter=100, multi_class='auto',
verbose=0, warm_start=False, n_jobs=None, l1_ratio=None)
    Initialize self. See help(type(self)) for accurate signature.

```

```

fit(self, X, y, sample_weight=None)
    Fit the model according to the given training data.

```

Parameters

-----

```

X : {array-like, sparse matrix} of shape (n_samples,
n_features)
    Training vector, where `n_samples` is the number of
samples and
    `n_features` is the number of features.

```

```

y : array-like of shape (n_samples,)
    Target vector relative to X.

```

```

sample_weight : array-like of shape (n_samples,) default=None
    Array of weights that are assigned to individual samples.
    If not provided, then each sample is given unit weight.

```

```

.. versionadded:: 0.17
    *sample_weight* support to LogisticRegression.

```

Returns

-----

```

self
    Fitted estimator.

```

## Notes

-----

The SAGA solver supports both float64 and float32 bit arrays.

`predict_log_proba(self, X)`

Predict logarithm of probability estimates.

The returned estimates for all classes are ordered by the label of classes.

## Parameters

-----

`X` : array-like of shape (n\_samples, n\_features)

Vector to be scored, where `n\_samples` is the number of samples and

`n\_features` is the number of features.

## Returns

-----

`T` : array-like of shape (n\_samples, n\_classes)

Returns the log-probability of the sample for each class

in the

model, where classes are ordered as they are in  
`self.classes\_`.

`predict_proba(self, X)`

Probability estimates.

The returned estimates for all classes are ordered by the label of classes.

For a multi\_class problem, if multi\_class is set to be "multinomial"

the softmax function is used to find the predicted probability of each class.

Else use a one-vs-rest approach, i.e calculate the probability of each class assuming it to be positive using the logistic function.

and normalize these values across all the classes.

## Parameters

-----

`X` : array-like of shape (n\_samples, n\_features)

Vector to be scored, where `n\_samples` is the number of samples and

`n\_features` is the number of features.

```

    Returns
    -----
    T : array-like of shape (n_samples, n_classes)
        Returns the probability of the sample for each class in
the model,
        where classes are ordered as they are in
``self.classes_``.

    set_fit_request(self:
sklearn.linear_model._logistic.LogisticRegression, *, sample_weight:
Union[bool, NoneType, str] = '$UNCHANGED$') ->
sklearn.linear_model._logistic.LogisticRegression
        Request metadata passed to the ``fit`` method.

        Note that this method is only relevant if
        ``enable_metadata_routing=True``
(see :func:`sklearn.set_config`).
        Please see :ref:`User Guide <metadata_routing>` on how the
routing
        mechanism works.

        The options for each parameter are:

        - ``True``: metadata is requested, and passed to ``fit`` if
provided. The request is ignored if metadata is not provided.

        - ``False``: metadata is not requested and the meta-estimator
will not pass it to ``fit``.

        - ``None``: metadata is not requested, and the meta-estimator
will raise an error if the user provides it.

        - ``str``: metadata should be passed to the meta-estimator
with this given alias instead of the original name.

        The default (``sklearn.utils.metadata_routing.UNCHANGED``)
retains the
        existing request. This allows you to change the request for
some
        parameters and not others.

        .. versionadded:: 1.3

        .. note::
            This method is only relevant if this estimator is used as
a
            sub-estimator of a meta-estimator, e.g. used inside a
            :class:`~sklearn.pipeline.Pipeline`. Otherwise it has no
effect.

```

```

Parameters
-----
sample_weight : str, True, False, or None,
default=sklearn.utils.metadata_routing.UNCHANGED
Metadata routing for ``sample_weight`` parameter in
``fit``.

Returns
-----
self : object
    The updated object.

set_score_request(self:
sklearn.linear_model._logistic.LogisticRegression, *, sample_weight:
Union[bool, NoneType, str] = '$UNCHANGED$') ->
sklearn.linear_model._logistic.LogisticRegression
    Request metadata passed to the ``score`` method.

    Note that this method is only relevant if
    ``enable_metadata_routing=True``
    (see :func:`sklearn.set_config`).
    Please see :ref:`User Guide <metadata_routing>` on how the
    routing
    mechanism works.

    The options for each parameter are:

    - ``True``: metadata is requested, and passed to ``score`` if
    provided. The request is ignored if metadata is not provided.

    - ``False``: metadata is not requested and the meta-estimator
    will not pass it to ``score``.

    - ``None``: metadata is not requested, and the meta-estimator
    will raise an error if the user provides it.

    - ``str``: metadata should be passed to the meta-estimator
    with this given alias instead of the original name.

    The default (``sklearn.utils.metadata_routing.UNCHANGED``)
    retains the
    existing request. This allows you to change the request for
    some
    parameters and not others.

    .. versionadded:: 1.3

    .. note::
        This method is only relevant if this estimator is used as

```

```

a
|         sub-estimator of a meta-estimator, e.g. used inside a
|         :class:`~sklearn.pipeline.Pipeline`. Otherwise it has no
effect.
|
|         Parameters
|         -----
|         sample_weight : str, True, False, or None,
default=sklearn.utils.metadata_routing.UNCHANGED
|         Metadata routing for ``sample_weight`` parameter in
``score``.
|
|         Returns
|         -----
|         self : object
|         The updated object.
|
|-----|
| Data and other attributes defined here:
|
| __annotations__ = {'_parameter_constraints': <class 'dict'>}
|
|-----|
| Methods inherited from
sklearn.linear_model._base.LinearClassifierMixin:
|
| decision_function(self, X)
|     Predict confidence scores for samples.
|
|     The confidence score for a sample is proportional to the
signed
|     distance of that sample to the hyperplane.
|
|     Parameters
|     -----
|     X : {array-like, sparse matrix} of shape (n_samples,
n_features)
|     The data matrix for which we want to get the confidence
scores.
|
|     Returns
|     -----
|     scores : ndarray of shape (n_samples,) or (n_samples,
n_classes)
|     Confidence scores per ``(n_samples, n_classes)``
combination. In the
|     binary case, confidence score for `self.classes_[1]` where

```



```

>0 means
    this class would be predicted.

predict(self, X)
    Predict class labels for samples in X.

    Parameters
    -----
    X : {array-like, sparse matrix} of shape (n_samples,
n_features)
        The data matrix for which we want to get the predictions.

    Returns
    -----
    y_pred : ndarray of shape (n_samples,)
        Vector containing the class labels for each sample.

```

---

```

Methods inherited from sklearn.base.ClassifierMixin:

score(self, X, y, sample_weight=None)
    Return the mean accuracy on the given test data and labels.

    In multi-label classification, this is the subset accuracy
    which is a harsh metric since you require for each sample that
    each label set be correctly predicted.

    Parameters
    -----
    X : array-like of shape (n_samples, n_features)
        Test samples.

    y : array-like of shape (n_samples,) or (n_samples, n_outputs)
        True labels for `X`.

    sample_weight : array-like of shape (n_samples,), default=None
        Sample weights.

    Returns
    -----
    score : float
        Mean accuracy of ``self.predict(X)`` w.r.t. `y`.

```

---

```

Data descriptors inherited from sklearn.base.ClassifierMixin:

__dict__

```

dictionary for instance variables (if defined)

`__weakref__`

list of weak references to the object (if defined)

---

Methods inherited from `sklearn.linear_model._base.SparseCoefMixin`:

`densify(self)`

Convert coefficient matrix to dense array format.

Converts the ```coef_``` member (back) to a `numpy.ndarray`. This is the default format of ```coef_``` and is required for fitting, so calling this method is only required on models that have previously been

sparsified; otherwise, it is a no-op.

Returns

-----

`self`

Fitted estimator.

`sparsify(self)`

Convert coefficient matrix to sparse format.

Converts the ```coef_``` member to a `scipy.sparse` matrix, which for L1-regularized models can be much more memory- and storage-efficient than the usual `numpy.ndarray` representation.

The ```intercept_``` member is not converted.

Returns

-----

`self`

Fitted estimator.

Notes

-----

For non-sparse models, i.e. when there are not many zeros in ```coef_```, this may actually *increase* memory usage, so use this method with care. A rule of thumb is that the number of zero elements, which can

be computed with `` (coef\_ == 0).sum() ``, must be more than 50%  
for this  
to provide significant benefits.

After calling this method, further fitting with the  
partial\_fit  
method (if any) will not work until you call densify.

---

Methods inherited from sklearn.base.BaseEstimator:

\_\_getstate\_\_(self)  
Helper for pickle.

\_\_repr\_\_(self, N\_CHAR\_MAX=700)  
Return repr(self).

\_\_setstate\_\_(self, state)

\_\_sklearn\_clone\_\_(self)

get\_params(self, deep=True)  
Get parameters for this estimator.

Parameters

-----

deep : bool, default=True

If True, will return the parameters for this estimator and  
contained subobjects that are estimators.

Returns

-----

params : dict

Parameter names mapped to their values.

set\_params(self, \*\*params)  
Set the parameters of this estimator.

The method works on simple estimators as well as on nested  
objects

(such as :class:`~sklearn.pipeline.Pipeline`). The latter have  
parameters of the form ``<component>\_\_<parameter>`` so that

it's  
possible to update each component of a nested object.

Parameters

-----

\*\*params : dict

```

    Estimator parameters.

    Returns
    -----
    self : estimator instance
        Estimator instance.

```

---

```

    Methods inherited from
    sklearn.utils._metadata_requests._MetadataRequester:

    get_metadata_routing(self)
        Get metadata routing of this object.

        Please check :ref:`User Guide <metadata_routing>` on how the
routing mechanism works.

    Returns
    -----
    routing : MetadataRequest
        A :class:`~sklearn.utils.metadata_routing.MetadataRequest`
encapsulating routing information.

```

---

```

    Class methods inherited from
    sklearn.utils._metadata_requests._MetadataRequester:

    __init_subclass__(**kwargs) from builtins.type
        Set the ``set_{method}_request`` methods.

        This uses PEP-487 [1]_ to set the ``set_{method}_request``
methods. It looks for the information available in the set default values
which are set using ``__metadata_request__*`` class attributes, or
inferred from method signatures.

        The ``__metadata_request__*`` class attributes are used when a
method does not explicitly accept a metadata through its arguments or
if the developer would like to specify a request value for those
metadata which are different from the default ``None``.

```


## References

.. [1] <https://www.python.org/dev/peps/pep-0487>

```
fruit_label = 1
y_one_vs_all = y==fruit_label
y_one_vs_all = y_one_vs_all.astype(int)

clf = LogisticRegression(C=1000).fit (X, y_one_vs_all)
print ('Accuracy={}'.format (clf.score(X,y_one_vs_all)))
plt.figure()
plot_decision_boundary(clf, X, y_one_vs_all,
title=fruits_dict[fruit_label], precision = 0.01)

Accuracy=0.7288135593220338
Call prediction for all grid values (precision of drawing = 0.01,
you may configure to speed up e.g. precision=0.05)
```




## Sklearn multiclass classification

```
print ('set(y) = {}'.format (set(y)))
print ('X.shape= {} \ny.shape = {}'.format (X.shape, y.shape))
clf= LogisticRegression(C=1000, max_iter=2000).fit(X, y)
print ('Accuracy={}'.format (clf.score(X,y)))
clf.intercept_, clf.coef_

set(y) = {1, 2, 3, 4}
X.shape= (59, 2)
y.shape = (59,)
Accuracy=0.847457627118644

(array([-29.75018322,  51.86931743, -41.76757778,  19.64844358]),
 array([[ 8.61970755, -3.51258454],
        [-1.47288929, -6.27247773],
        [ 6.30769266,  0.33620857],
        [-13.45451092,  9.4488537 ]]))

plt.figure()
plot_multi_class_logistic_regression (X,y,dict_names=fruits_dict)
draw_linear_decision_boundaries_multiclass(clf,X)
```



## Iris dataset

```
from sklearn.datasets import load_iris
iris = load_iris()
X, y, labels = iris.data, iris.target, iris.target_names
print (labels)
clf= LogisticRegression(C=100, max_iter=2000).fit(X, y)
print ('Accuracy={} '.format (clf.score(X,y)))
clf.intercept_, clf.coef_

['setosa' 'versicolor' 'virginica']
Accuracy=0.98

(array([ 19.94628563,   5.23090235, -25.17718798]),
 array([[ -0.39380232,   3.41248293, -6.40562604, -3.50688137],
        [ 1.35772093,   0.44540853, -0.51275436, -4.41520581],
        [-0.96391861, -3.85789146,   6.9183804 ,   7.92208718]]))

features = ['sepal_length', 'sepal_width', 'petal_length',
'petal_width']
dict_names = {i:v for i,v in enumerate(labels)}

X2= X[:, :2]
clf= LogisticRegression(C=100).fit(X2, y)
print ('Accuracy={} '.format (clf.score(X2,y)))
print ('clf.intercept_={}, \nclf.coef_=\n{}'.format(clf.intercept_,
clf.coef_))
plt.figure()
plot_multi_class_logistic_regression (X2,y, dict_names = dict_names)
draw_linear_decision_boundaries_multiclass(clf,X2)

Accuracy=0.8333333333333334
clf.intercept_=[ 25.20698869 -6.10712004 -19.09986865],
clf.coef_ =
[[-9.33565746   8.18421204]
 [ 3.72540037 -4.30414546]
 [ 5.61025709 -3.88006658]]
```

